



# Introduction to Machine Learning

## Introduction

林彥宇 教授

Yen-Yu Lin, Professor


國立陽明交通大學 資訊工程學系

Computer Science, National Yang Ming Chiao Tung University

Some slides are modified from Prof. Sheng-Jyh Wang,  
Prof. Hwang-Tzong Chen, and Prof. Yung-Yu Chuang

# Pattern recognition and machine learning

- Pattern recognition is the automated recognition of patterns and regularities in data
  - Discover pattern regularities
  - Take actions, such as classification or regression, with regularities
- Data: A set of hand-written digits and the class ground truth



- Computer algorithm: It extracts features from each image, analyze the patterns and regularities in data
- Model: Given a new hand-written digit, predict its class label

# Pattern recognition and machine learning

- Machine learning: to design and develop algorithms that allow computers to predict data based on empirical data
  - Try to explore certain patterns or regularities
  - Learn models from the given data
  - Based on the given data, the learner produces a useful output in new cases
- Machine learning is one approach to pattern recognition, while other approaches include hand-crafted (not learned) rules or heuristics
- Machine learning  $\subset$  Pattern recognition

# Applications of machine learning


- Computer vision
- Speech recognition
- Information retrieval
- Natural language processing
- Robotics
- Bioinformatics
- Data mining
- Finance
- ...

# Problem definition of a machine learning task

- **Training data**
  - A set of  $N$  **training data**  $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ , sometimes together with their **target vectors**  $\{t_1, t_2, \dots, t_N\}$
- **Feature extraction**
  - Original input variables are usually transformed into some new space of variables, where the problem can be better handled
- **Model learning**
  - We learn a proper model for the problem
- **Generalization or testing**
  - To correctly predict new examples (**testing data**) that differ from those used for training


# Cat image classification: Training data

- Collect a set of **training data** with **target vectors**



# Cat image classification: Feature extraction

- Feature extraction is crucial
  - Need to take feature variations into account



viewpoint variations



Illumination variations



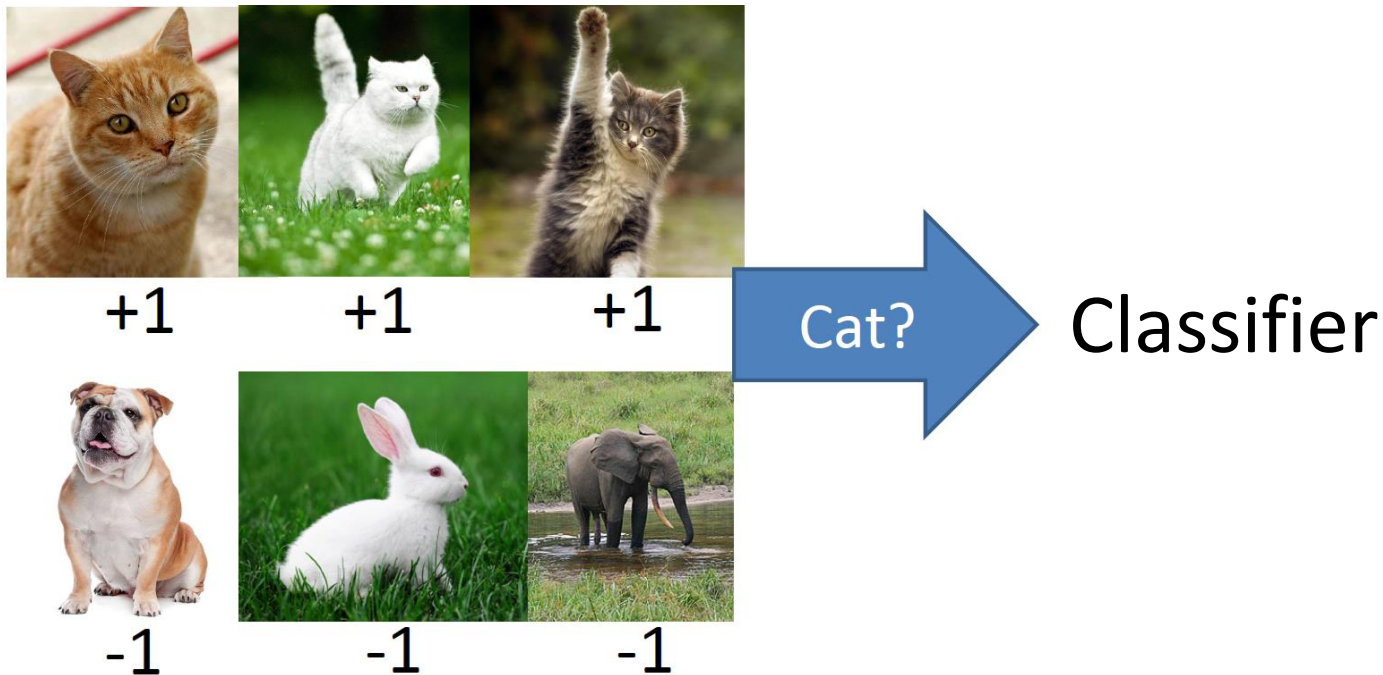
background variations



pose variations

# Cat image classification: Model learning

- Based on the given training data and the extracted features, we learn a classifier





# Cat image classification: Testing

- Apply the learned **classifier** to the testing images




# Cat image classification: Testing

- Apply the learned **classifier** to the testing images and make **prediction**




# Regression



資料日期: 2014-09-29 ~ 2015-10-06

累計成交金額: 229,767.44 億



● 加權指數

■ 成交金額 (億)



estimated TAIEX on 11/5

Taiwan Capitalization  
Weighted Stock Index


# Supervised vs. Unsupervised learning

- **Supervised learning**: the training data comprises examples of the input vectors **along with** their corresponding target vectors
  - **Classification**: assign each input vector to one of a finite number of discrete categories
  - **Regression**: assign each input vector to one or more continuous variables
  - Methods: linear regression, linear classification, neural networks, support vector machine, ensemble learning, dimensionality reduction, deep learning, ...

# Good vs. bad features for classification



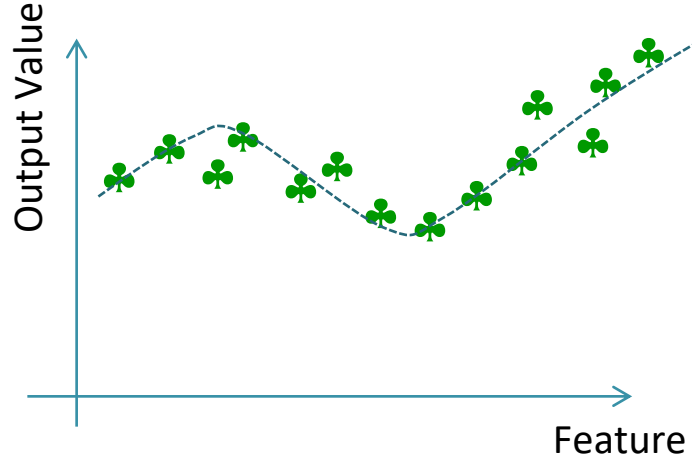

good features




bad features



# Good vs. bad features for regression



good feature




bad feature

# Supervised vs. Unsupervised learning

- **Unsupervised learning**: the training data consist of a set of input vectors  $x$  **without** any corresponding target values
  - **Clustering**: to discover groups of similar examples within the data
  - **Density estimation**: to determine the distribution of data within the input space
  - **Dimensionality reduction**: to project the data from a high-dimensional space down to a low-dimensional space
  - **Data generation**: to synthesize new data with some particular conditions

# Unsupervised learning for clustering

- **Clustering**: To group a set of data in such a way that data points in the same group, called a **cluster**, are more similar to each other than to those in other clusters

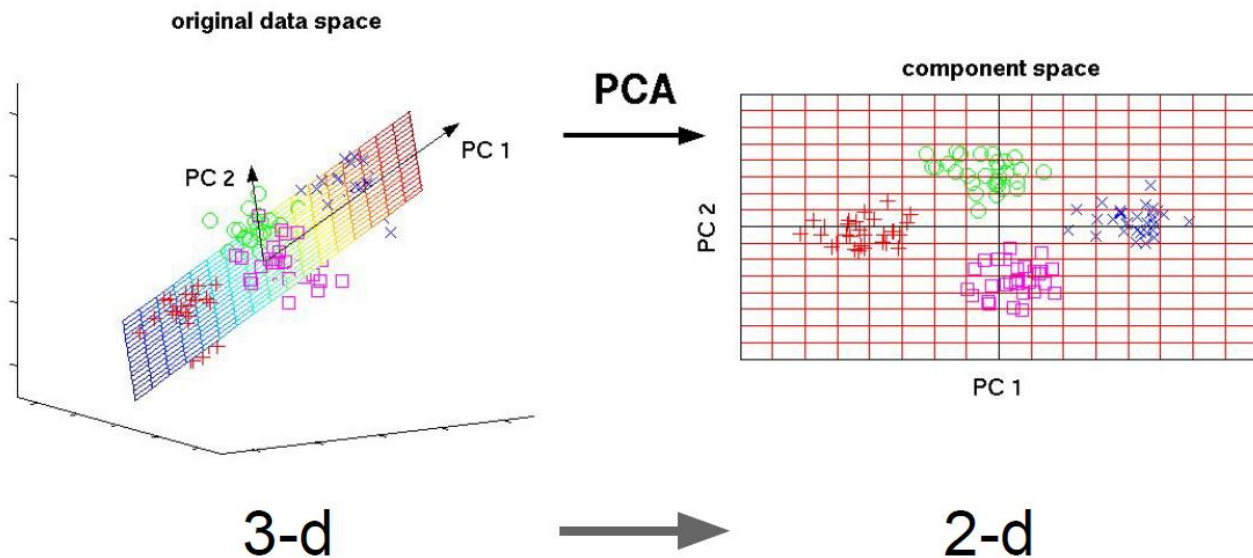


***k*-mean clustering**



# Unsupervised learning for dimensionality reduction


- **Dimensionality reduction**: To project data from a high-dimensional space to a low-dimensional one



**PCA: Principal  
component  
analysis**

# Unsupervised learning for density estimation


- **Density estimation:** Based on given data, estimate the underlying probability density function



**kernel density  
estimation (KDE)**

# Unsupervised learning for data generation

- Given a set of natural images, we try to generate new images that look natural and photorealistic



Training data  $\sim p_{\text{data}}(x)$



Generated samples  $\sim p_{\text{model}}(x)$

**Generative Adversarial Networks (GAN):** Given a set of images, generate new images from the same distributions

# Applications of data generation


- Face synthesis



Tero Karras et al. "Progressive Growing of GANs for Improved Quality, Stability, and Variation"

# Polynomial curve fitting: Problem definition

- Training data (observations)
  - 10 blue circles, each of which has
    - ◆ One-dimensional input
    - ◆ One target output
- Green curve  $\sin(2\pi x)$  is the function used to generate these data, which is **unknown**
- Each point is sampled from the function with a random Gaussian noise
- **Goal of curve fitting:** To exploit the training data to discover the underlying function so that we can make predictions of the value  $\hat{t}$  for some new input  $\hat{x}$



# Polynomial curve fitting: Choose a fitting function

- Fit the data using a polynomial function of the form:

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

- This function is **parametrized** by  $\mathbf{w}$
- $w_0$  is the bias term
- Its input is a data point, while the output is estimated target
- $M$  is the **order** of the polynomial function



# Polynomial curve fitting: Error function


- An error function (objective function) is used to determine the parameters

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

- In this case, we minimize the **sum-of-squares error**

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$





- Differentiable
- Closed form solution



# Polynomial curve fitting: **Model selection**

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$


- Models with different values of **hyperparameter**  $M$



- Model selection:** To choose a proper value of  $M$



# Polynomial curve fitting: Model selection




- **Under-fitting:**  $M = 0$  or  $M = 1$ 
  - The constant or first order polynomial gives poor fit due to insufficient flexibility
- The third order polynomial gives the best fit
- **Over-fitting:**  $M = 9$ 
  - All training points are perfectly fitted
  - Poor representation of the green curve
  - The generalization is poor



# Polynomial curve fitting: Generalization


- Suppose we are given a set of training data and a separate set of 100 test data
- Evaluate the generalization for each choice of  $M$  via **root-mean-square (RMS) error**

$$E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$$



	$M = 0$	$M = 1$	$M = 6$	$M = 9$
$w_0^*$	0.19	0.82	0.31	0.35
$w_1^*$		-1.27	7.99	232.37
$w_2^*$			-25.43	-5321.83
$w_3^*$			17.37	48568.31
$w_4^*$				-231639.30
$w_5^*$				640042.26
$w_6^*$				-1061800.52
$w_7^*$				1042400.18
$w_8^*$				-557682.99
$w_9^*$				125201.43

# Polynomial curve fitting: Generalization






- Small values of  $M$  give relatively large values of training and test errors
- When  $M$  is between 3 and 8, reasonable representations are obtained
- For  $M=9$ , the training error goes to zero, but the test error increases significantly



# Polynomial curve fitting: Data size vs. Over-fitting

$M = 9$



- Over-fitting becomes less severe as the data size increases
- In general, the number of data points should be no less than some multiple (say 5 or 10) of the number of adaptive parameters in the model
- **Regularization** is often used to control the over-fitting phenomenon

# Polynomial curve fitting: Regularization



- **Regularization**: Add a penalty term to the error function to discourage the coefficients from reaching large values

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$


where  $\|\mathbf{w}\|^2 = \mathbf{w}^T \mathbf{w} = \omega_0^2 + \omega_1^2 + \cdots + \omega_M^2$

- The coefficient  $\omega_0$  is usually omitted
- This kind of techniques is called **shrinkage** methods in the statistics literature
- A quadratic regularizer is called **ridge regression**
- In neural networks, this approach is known as **weight decay**

# Polynomial curve fitting: Regularization




	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
$w_0^*$	0.35	0.35	0.13
$w_1^*$	232.37	4.74	-0.05
$w_2^*$	-5321.83	-0.77	-0.06
$w_3^*$	48568.31	-31.97	-0.05
$w_4^*$	-231639.30	-3.89	-0.03
$w_5^*$	640042.26	55.28	-0.02
$w_6^*$	-1061800.52	41.32	-0.01
$w_7^*$	1042400.18	-45.95	-0.00
$w_8^*$	-557682.99	-91.53	0.00
$w_9^*$	125201.43	72.68	0.01



# Probability theory

- We need to handle data **uncertainties**, which result from
  - Noise on measurement
  - Finite size of data sets
- Probability theory provides a consistent framework to manipulate uncertainties, and hence is essential to pattern recognition research

# A toy examples




- Two boxes:  $r$  (red box) and  $b$  (blue box)
- Two types of fruits:  $a$  (apple) and  $o$  (orange)
- A **trial**: Randomly selecting a box from which we randomly picking a fruit
- Introduce one variable  $B$  for box and one variable  $F$  for fruit
- Many trials: Repeat the process many times
- Question 1: What is the probability that an apple is picked
  - **Marginal probability**
- Question 2: Given that we have picked an orange, what is the probability that the box we chose was the blue one?
  - **Conditional probability**




# Probability theory: A two-variable case

- Two random variables:  $X$  and  $Y$
- Each variable has a set of discrete states
  - $X$  can take any value  $x_i$  where  $i = 1, 2, \dots, M$
  - $Y$  can take any value  $y_j$  where  $j = 1, 2, \dots, L$
- $N$  trails where both variables  $X$  and  $Y$  are sampled
- Some notations
  - Let the number of trails where  $X = x_i$  and  $Y = y_j$  be  $n_{ij}$
  - Let the number of trails where  $X$  takes value  $x_i$  be  $c_i$
  - Let the number of trails where  $Y$  takes value  $y_j$  be  $r_j$



# Joint, marginal, and conditional probabilities



- The probability that  $X$  takes value  $x_i$  and  $Y$  takes value  $y_j$  is called **joint probability**
- It is defined by the fraction of points (trials) falling in the cell  $i,j$

$$p(X = x_i, Y = y_j) = \frac{n_{ij}}{N}$$

# Joint, **marginal**, and conditional probabilities


- The probability that  $X$  takes value  $x_i$  irrespective of the value of  $Y$  is called **marginal probability** and is written as  $p(X = x_i)$
- It is defined by the fraction of the number of points that fall in column  $i$ , namely

$$p(X = x_i) = \frac{c_i}{N}$$

- With the joint probability and  $c_i = \sum_j n_{ij}$ , we have

$$p(X = x_i) = \sum_{j=1}^L p(X = x_i, Y = y_j)$$

- The **sum rule**




# Joint, marginal, and conditional probabilities

- If we consider only those cases where  $X$  takes value  $x_i$ , the fraction of those cases where  $Y = y_j$  is written as  $p(Y = y_j | X = x_i)$ . It is called conditional probability

- It is defined by

$$p(Y = y_j | X = x_i) = \frac{n_{ij}}{c_i}$$



- Relationships among joint, marginal, and conditional probabilities:

$$\begin{aligned} p(X = x_i, Y = y_j) &= \frac{n_{ij}}{N} = \frac{n_{ij}}{c_i} \cdot \frac{c_i}{N} \\ &= p(Y = y_j | X = x_i) p(X = x_i) \end{aligned}$$

- The product rule





# Joint, marginal, and conditional probabilities

sum rule

$$p(X) = \sum_Y p(X, Y)$$

product rule

$$p(X, Y) = p(Y|X)p(X)$$



# Bayes' theorem

- By using the product rule and the symmetry property  $p(X, Y) = p(Y, X)$ , we have

$$p(Y|X) = \frac{p(X|Y)p(Y)}{p(X)}$$



# Probability with continuous variables

- The **probability density**  $p(x)$  over a continuous variable  $x$  must satisfy the two conditions:

$$p(x) \geq 0$$

$$\int_{-\infty}^{\infty} p(x) dx = 1$$

- **Nonnegative**: Probabilities are nonnegative
  - **Sum-to-1**: The value of  $x$  must lie somewhere on the real axis
- The **cumulative distribution function** defines the probability that  $x$  lies in the interval  $(-\infty, z)$  via

$$P(z) = \int_{-\infty}^z p(x) dx$$



# Sum rule and product rule

- Sum rule in discrete cases

$$p(X) = \sum_Y p(X, Y)$$

- Sum rule in continuous cases

$$p(x) = \int p(x, y) \, dy$$

- Product rule in discrete cases

$$p(X, Y) = p(Y|X)p(X)$$

- Product rule in continuous cases

$$p(x, y) = p(y|x)p(x)$$



# Expectations and covariances

- The **average value of some function**  $f(x)$  under a probability distribution  $p(x)$  is called the **expectation** of  $f(x)$
- For a discrete distribution, the expectation of  $f(x)$  is

$$\mathbb{E}[f] = \sum_x p(x) f(x)$$

- For a continuous probability, the expectation of  $f(x)$  is

$$\mathbb{E}[f] = \int p(x) f(x) \, dx$$

# Expectations and covariances

- The **variance** of  $f(x)$  under a probability distribution  $p(x)$  is

$$\text{var}[f] = \mathbb{E} [(f(x) - \mathbb{E}[f(x)])^2]$$

- It is a measure of how much variability there is in  $f(x)$  around its mean  $\mathbb{E}[f(x)]$
- For two random variables  $x$  and  $y$ , the **covariance** is defined by

$$\begin{aligned}\text{COV}[x, y] &= \mathbb{E}_{x,y} [\{x - \mathbb{E}[x]\} \{y - \mathbb{E}[y]\}] \\ &= \mathbb{E}_{x,y} [xy] - \mathbb{E}[x]\mathbb{E}[y]\end{aligned}$$


- It expresses the extent to which  $x$  and  $y$  vary together.

# Gaussian distribution

- For a single continuous variable, the **Gaussian** or **normal distribution** is defined by

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp \left\{ -\frac{1}{2\sigma^2} (x - \mu)^2 \right\}$$

which is specified by two parameters: **mean  $\mu$**  and **variance  $\sigma^2$**



$$\mathcal{N}(x|\mu, \sigma^2) > 0$$

$$\int_{-\infty}^{\infty} \mathcal{N}(x|\mu, \sigma^2) dx = 1$$



# Mean and variance of a Gaussian distribution

- The **average value** of a random variable  $x$  whose distribution is Gaussian

$$\mathbb{E}[x] = \int_{-\infty}^{\infty} \mathcal{N}(x|\mu, \sigma^2) x \, dx = \mu$$

- The **second order moment** of variable  $x$

$$\mathbb{E}[x^2] = \int_{-\infty}^{\infty} \mathcal{N}(x|\mu, \sigma^2) x^2 \, dx = \mu^2 + \sigma^2$$

- The **variance** of variable  $x$

$$\text{var}[x] = \mathbb{E}[x^2] - \mathbb{E}[x]^2 = \sigma^2$$

# Multivariate Gaussian

- The **multivariate Gaussian distribution** defined over a  $D$ -dimensional vector  $\mathbf{x}$  of continuous variables:

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\boldsymbol{\Sigma}|^{1/2}} \exp \left\{ -\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu}) \right\}$$

where  $D \times D$  matrix is called the **co-variance** matrix while  $|\boldsymbol{\Sigma}|$  denotes the determinant of  $\boldsymbol{\Sigma}$

# Bayes' theorem for polynomial curve fitting

- Recall the curve fitting problem
  - Given a set of  $N$  observations  $D = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$  and their target values  $\{t_1, t_2, \dots, t_N\}$
  - Polynomial curve fitting: Determine the values of  $\mathbf{w}$

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

- **Prior probability**  $p(\mathbf{w})$ : Express our assumption about  $\mathbf{w}$  **before** observing any data
- **Likelihood function**  $p(D|\mathbf{w})$ : Express how probable the observed data  $D$  is under  $\mathbf{w}$ . It is evaluated **after** the observations  $D$  are given

# Bayes' theorem for polynomial curve fitting

- Bayes' theorem takes the form

$$p(\mathbf{w}|\mathcal{D}) = \frac{p(\mathcal{D}|\mathbf{w})p(\mathbf{w})}{p(\mathcal{D})}$$

which allows us to evaluate the uncertainty after we have observations  $\mathcal{D}$

- $p(\mathcal{D})$  is the normalization constant. Thus, we have

$$\text{posterior} \propto \text{likelihood} \times \text{prior}$$

# Determining Gaussian parameters by maximum likelihood

- Given a set of  $N$  observations:  $\mathbf{x} = (x_1, \dots, x_N)$
- Assume these observations are sampled from a Gaussian distribution with mean  $\mu$  and variance  $\sigma^2$  (unknown)
- Our goal is to determine  $\mu$  and  $\sigma^2$  based on the observations
- We assume that data are sampled independently from the same distribution, namely **independent and identically distributed**, or **i.i.d.** for short



# Determining Gaussian parameters by maximum likelihood

- Since the data are i.i.d., the **likelihood function** of data given mean  $\mu$  and variance  $\sigma^2$  is

$$p(\mathbf{x}|\mu, \sigma^2) = \prod_{n=1}^N \mathcal{N}(x_n|\mu, \sigma^2)$$

- The **log likelihood function** is

$$\ln p(\mathbf{x}|\mu, \sigma^2) = -\frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2 - \frac{N}{2} \ln \sigma^2 - \frac{N}{2} \ln(2\pi)$$

- **Maximum likelihood** solution:

$$\mu_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N x_n \quad \sigma_{\text{ML}}^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \mu_{\text{ML}})^2$$

# Probabilistic perspective of polynomial curve fitting

- Given  $N$  data for regression:  $\mathbf{x} = (x_1, \dots, x_N)^T$  &  $\mathbf{t} = (t_1, \dots, t_N)^T$ 
  - Fit the data using a polynomial function of the form:

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$


- This function is parametrized by  $\mathbf{w}$
- Given the value of  $x$ , we **assume** the corresponding value of  $t$  has a **Gaussian distribution** with a mean equal to  $y(x, \mathbf{w})$

$$p(t|x, \mathbf{w}, \beta) = \mathcal{N}(t|y(x, \mathbf{w}), \beta^{-1})$$

where  $\beta^{-1}$  is the variance  $\sigma^2$  ( $\beta$  is called precision)

# Probabilistic perspective of polynomial curve fitting

- The Gaussian conditional distribution for  $t$  given  $x$



- If data are i.i.d., the likelihood function is

$$p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n|y(x_n, \mathbf{w}), \beta^{-1})$$

# Maximum likelihood solution

- The log likelihood function

$$\ln p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta) = -\frac{\beta}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi)$$

- Maximum likelihood (ML) solution for determining  $\mathbf{w}$  and  $\beta$ 
  - Compute the gradient of the log likelihood function w.r.t.  $\mathbf{w}$ . And set it to 0. We can get  $\mathbf{w}_{\text{ML}}$ .

$$\sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$


- By setting the gradient of the log likelihood function w.r.t.  $\beta$  to 0,  $\beta_{\text{ML}}$  is obtained by solving

$$\frac{1}{\beta_{\text{ML}}} = \frac{1}{N} \sum_{n=1}^N \{y(x_n, \mathbf{w}_{\text{ML}}) - t_n\}^2$$

# Maximum likelihood solution

- After determining the values of  $\mathbf{w}_{\text{ML}}$  and  $\beta_{\text{ML}}$ , we can make predictions for a new value of  $x$

$$p(t|x, \mathbf{w}_{\text{ML}}, \beta_{\text{ML}}) = \mathcal{N}(t|y(x, \mathbf{w}_{\text{ML}}), \beta_{\text{ML}}^{-1})$$



# Maximum a posterior (MAP) solution

- While ML solution is obtained by maximizing the likelihood, MAP solution is by maximizing the posterior
- Recall  $\text{posterior} \propto \text{likelihood} \times \text{prior}$
- Introduce a prior distribution over the curve parameters  $\mathbf{w}$

$$p(\mathbf{w}|\alpha) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^{-1}\mathbf{I}) = \left(\frac{\alpha}{2\pi}\right)^{(M+1)/2} \exp\left\{-\frac{\alpha}{2}\mathbf{w}^T\mathbf{w}\right\}$$

- $M$  is the order of the polynomial
  - $\alpha$  is a hyperparameter
- The posterior distribution for  $\mathbf{w}$

$$p(\mathbf{w}|\mathbf{x}, \mathbf{t}, \alpha, \beta) \propto p(\mathbf{t}|\mathbf{x}, \mathbf{w}, \beta)p(\mathbf{w}|\alpha).$$

# Maximum a posterior (MAP) solution

- The MAP solution,  $\mathbf{w}_{\text{MAP}}$  and  $\beta_{\text{MAP}}$ , is obtained by maximizing the posterior function, or equivalently by minimizing

$$\frac{\beta}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\alpha}{2} \mathbf{w}^T \mathbf{w}.$$

# Bayesian curve fitting

- We make a **point** estimation of  $\mathbf{w}$  no matter in ML and MAP solutions
- In a full Bayesian approach, we integrate over all possible values of  $\mathbf{w}$  for regression, i.e.,


$$p(t|x, \mathbf{x}, \mathbf{t}) = \int p(t|x, \mathbf{w})p(\mathbf{w}|\mathbf{x}, \mathbf{t}) d\mathbf{w} = \mathcal{N}(t|m(x), s^2(x))$$

$$m(x) = \beta \phi(x)^T \mathbf{S} \sum_{n=1}^N \phi(x_n) t_n$$

$$s^2(x) = \beta^{-1} + \phi(x)^T \mathbf{S} \phi(x).$$

$$\phi(x_n) = (x_n^0, \dots, x_n^M)^T$$

$$\mathbf{S}^{-1} = \alpha \mathbf{I} + \beta \sum_{n=1}^N \phi(x_n) \phi(x_n)^T$$





# Probabilistic polynomial curve fitting

- Given the assumption  $p(t|x, \mathbf{w}, \beta) = \mathcal{N}(t|y(x, \mathbf{w}), \beta^{-1})$ 
  - **ML solution**: Find  $\mathbf{w}$  that maximizes the likelihood function

$$p(t | x, D) = p(t | x, \mathbf{w}_{\text{ML}}, \beta^{-1})$$

- **MAP solution**: Find  $\mathbf{w}$  that maximizes the posterior probability

$$p(t | x, D) = p(t | x, \mathbf{w}_{\text{MAP}}, \beta^{-1})$$

- **Bayesian solution**: Integrate over  $\mathbf{w}$

$$p(t | x, D) = p(t|x, \mathbf{x}, \mathbf{t}) = \int p(t|x, \mathbf{w})p(\mathbf{w}|\mathbf{x}, \mathbf{t}) d\mathbf{w}$$



# Model selection

- Hyperparameters, such as  $M$  in polynomial curve fitting, control the model behavior complexity

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

- Model selection: determine the values of hyperparameters that achieve the **best predictive performance on new (testing) data**
- Idea: split training data into a training set and **a validation set**
  - Training set: Used to learn the model with particular hyperparameters values
  - Validation set: Used to evaluate the performance of the learned model


# Model selection

- About the size of the validation set
  - A large validation set: Less training data for model learning
  - A small validation set: Less reliable performance evaluation

# Model selection via cross validation

- $S$ -fold cross-validation

- Partition training data into  $S$  equal-sized groups
- $S-1$  groups are used to train the model that is evaluated on the remaining group
- Repeat the procedure for all  $S$  possible runs
- Average the performance



# Drawbacks of model selection

- If training data are limited, a large value of  $S$  is appropriate
- At the extreme, setting  $S=N$  (number of training data), it gives the **leave-one-out** technique
- Some drawbacks
  - The number of training runs increases by a factor of  $S$
  - The number of hyperparameter value combinations increases exponentially

# Summary

- Polynomial curve fitting for regression
  - Fitting by minimizing the sum-of-squares error

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

- Regularization for alleviating overfitting

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- Probability density
  - Expectation, variance, and covariance
  - Gaussian distribution

# Summary

- Bayes' theorem

$$\text{posterior} \propto \text{likelihood} \times \text{prior}$$

- When applying Bayes' theorem to polynomial curve fitting,
  - ML solution: Find  $\mathbf{w}$  that maximizes the likelihood function
  - MAP solution: Find  $\mathbf{w}$  that maximizes the posterior probability
  - Bayesian solution: Integrate over  $\mathbf{w}$
- Model selection by cross-validation

# References

- Chapters 1.1, 1.2, 1.3, and 1.4 in the PRML textbook



# Thank You for Your Attention!

THANK YOU FOR YOUR ATTENTION!

**Yen-Yu Lin (林彥宇)**

Email: [lin@cs.nctu.edu.tw](mailto:lin@cs.nctu.edu.tw)

URL: <https://www.cs.nycu.edu.tw/members/detail/lin>